Toxicity Detection in Multiplayer Online Games

Marcus Märtens*, Siqi Shen†, Alexandru Iosup‡ and Fernando Kuipers*

*Network Architectures and Services Group, Delft University of Technology, Delft, The Netherlands
†Parallel and Distributed Processing Laboratory, National University of Defense Technology, China
†College of Computer, National University of Defense Technology, China
‡Parallel and Distributed Systems Group, Delft University of Technology, Delft, The Netherlands
Email: {m.maertens, f.a.kuipers, a.iosup}@tudelft.nl, shensiqi@nudt.edu.cn

Abstract—**Social interactions in multiplayer online games are an essential feature for a growing number of players world-wide. However, this interaction between the players might lead to the emergence of undesired and unintended behavior, particularly if the game is designed to be highly competitive. Communication channels might be abused to harass and verbally assault other players, which negates the very purpose of entertainment games by creating a toxic player-community. By using a novel natural language processing framework, we detect profanity in chat-logs of a popular Multiplayer Online Battle Arena (MOBA) game and develop a method to classify toxic remarks. We show how toxicity is non-trivially linked to game success.**

I. Introduction

Multiplayer Online Battle Arena (MOBA) games have been growing increasingly popular and captivate their player base in virtue of complex game mechanics and competitive nature. Riot's League of Legends claims to have over 67M monthly active players[1] and grosses over 1 billion US dollars of revenue yearly.[2] With 18M US dollars, one of the largest price pools in the history of eSports for a single tournament was crowdfunded almost entirely by the player base of Valve's Dota 2.[3]

MOBAs are played in independent n vs n matches, typically with $n = 5$, in which the players of each team need to closely cooperate to penetrate the other team's defences and obtain victory. Players who refuse to cooperate and act without considering their own team are easy targets and get killed more frequently, which diminishes the team's chances. Together with the intricate and sometimes counter-intuitive strategic nature of MOBAs, this gives rise to conflict within the teams. Triggered by game events like kills or just simple mistakes, players begin to turn sour. The communication channels that were meant to coordinate the team effort can then be used to verbally assault other players, often by using profane terms and heavy insults.

Possible consequences are resigned players, whom might no longer be interested in competing for the win. But even if the match is won eventually, players could still feel offended, abused and might regret their decision to play the game in general. In this way, the *mood* of a communication could qualify as a social Quality of Experience (QoE) metric [1].

Collecting bad game experiences like this is harmful for the community, as it can bias a player's attitude towards engaging in cooperation even when confronted with fresh opponents and new teammates in later matches. The perceived hostility in a player community is frequently referred to as *toxicity*. Toxicity imposes a serious challenge for game designers, as it may chase active regular players away. It might also prevent new players from joining the game, because a toxic base appears as unfriendly and hostile to newcomers.

The main contribution of this work is to devise an annotation system for chats of multiplayer online games that can be used for detecting toxicity (Section III). We apply the system to a large dataset (Section II) collected from a representative game of the MOBA genre and propose a method based on machine learning that uses the annotation system to predict the outcome of ongoing matches (Section IV). We end with related work (Section V) and conclusions (Section VI).

II. Data

A. Data sources

All data used in this work are based on one of the ancestors of all MOBA games: Defense of the Ancients (DotA).[4] This game started as a custom map for the real-time strategy game Warcraft III but soon became so popular that community platforms emerged that allowed for players to register, get profiled and being matched up against each other based on their skill. One of these platforms was DotAlicious, from which we crawled our data.

The website of DotAlicious is no longer available online, as DotA has been substituted by newer MOBAs like League of Legends, Heroes of the Storm or Dota II. The core game principles have not been changed by much by DotA's successors, but the accessibility of replays, chat-logs and player-related information for them is more limited due to several privacy concerns of the developing companies. Also, alternative means of information exchange, like protected voice-chats, make it more difficult to obtain a record of comprehensive inter-team communication. Hence, we believe that our data from DotA are suitable for our purpose while still being representative for the game genre in general. Additionally, it allows us to study toxicity without harming a *live* community.

---

[1] http://goo.gl/LHd8WJ (www.riotgames.com, Sep. 2015)
[2] http://goo.gl/bBKggU (gamasutra.com, Sep. 2015)
[3] http://goo.gl/FuFK6u (www.theguardian.com, Sep. 2015)

[4] http://www.playdota.com

Fig. 1. Distribution of match duration in the DotAlicious dataset.

B. Data cleansing and match outcome

Our DotAlicious dataset consists of replays from 12923 matches, spanning the time between the 2nd and the 6th of February 2012.

The duration of matches in the dataset is distributed bimodally, indicating that a small fraction of the matches ended prematurely. We used information from the hosting-bot of DotAlicious to determine matches that resulted in a draw or were canceled by the players early on. In total, out of 12923 matches, 1653 were aborted before game start, 706 were canceled after game start and 241 resulted in a draw by mutual player agreement (see Figure 1).

For the remaining matches, there are two possible outcomes: either one team destroys the other team's main structure (victory condition) or all players of one team forfeit, which results in a collective surrender (loss condition). We have identified 10305 matches with a well-defined winning team, of which 6082 matches ended by the victory condition and 4223 matches by surrender. 18 matches needed to be excluded as their outcomes were unclear.

III. GAME COMMUNICATION MODELLING

A. Annotation system design

For all matches, we extracted all chat-lines used by the players and applied a tokenization based on simple whitespace splitting. Symbols like "!" or "?" remained part of the words, as long as they were not separated by whitespaces. The case of the letters was unchanged to analyse the use of capitalization as a stylistic figure (shouting).

Overall, the language used is extremely abbreviated, elliptical, full of spelling-errors and barely following grammatical structures. Consequently, standard techniques from Natural Language Processing (NLP) like part-of-speech recognition, spelling-correction and language detection were either not applicable or performed poorly. On the other hand, we observed little variety in the topic of the chat, resulting in a rather restricted and repetitive set of vocabulary. We thus devised a novel annotation system to classify the most frequent words together with their miss-spelled variants.

The most dominant language in the corpus is English, which is used as a *pidgin language* for non-native speakers to communicate with each other. To classify the most frequently used words in this work, we do not consider words from any other language. Consequently, non English words will be either "unannotated" or classified as "non-latin" (for example in the case of Chinese, which is easy to detect).

To classify the semantics of a word, we apply sets of simple rules to them. There are three different classes of rules that we use:

1) **pattern**: the word includes or starts with certain symbols,
2) **list**: the word is member of a pre-defined list, and
3) **letterset**: the set of letters of the word equals the set of letters of a word from a pre-defined list.

The letterset class is useful to capture unintentionally or intentionally misspelled words, if no meaningful recombinations of their letters (like anagrams) exist in the corpus. For example, the set of letters used to spell the word "noob"[5] is {"n","o","b"}, which is the same set as used for words like "NOOOOOOOOb", "boon", "nooobbbbb" or "noonb" which were actually used in the chats. In total, the letterset method allowed to capture 224 (case-sensitive) different ways of writing "noob" that were used in the dataset. On the other hand, no other meaningful English word that could be built using this set (for example "bonobo" or "bonbon") was found in the corpus. Also for other words than "noob", the amount of introduced *false positives* due to the letterset-method was negligible for our dataset.

The precise word-lists and patterns that we used are provided as supplements to this work online[6] as they are technical details. Nevertheless, Table I shows the rule classes used for each annotation category together with a short description, some examples, their precedence and their absolute prevalence in the text corpus.

The text-corpus consists of 7042112 words in total, of which 286654 are distinct. Each distinct word is checked against our rules and annotated accordingly. If no rules apply, the word is "unannotated". If multiple rules apply, we break the tie by choosing the category with the highest precedence. Considering the set of all distinct words in the corpus, our annotation system covers around 16% of them. However, many of the most-frequent words are annotated, so that on average over 60% of all (non-distinct) words used per match have an annotation.

B. Different chat-modes

Our data allows us to investigate two fundamentally different chat-modes for each match: in the all-chat, a player can broadcast a message to each other player that participates in the match. In the ally-chat, the message is only sent to

---

[5]"Noob" is a common insult in video games. It is derived from the word "newbie", which comes from "newcomer". Thus, it implies that someone has the lowest possible level of skill and knowledge of the game.

[6]https://www.nas.ewi.tudelft.nl/people/mmaertens/toxicity/supplement_toxicity.txt

| category	description	rules	examples	precedence	unique count
nonlatin	special character, foreign language	pattern	文章	500	20133
praise	acts of courtesy, kindness, sport spirit or gratitude	list	gj, gg, thx, hf	100	295
bad	profanity, swear words, inappropriate language	list, letterset	noob, idiot, f*	90	4881
laughter	acronyms expressing laughter	letterset	HAHAHAHA, lol, ROFL	60	2158
smiley	emoticons, symbols resembling faces or emotions	pattern, list	:D, ;), oO, -_-	50	1110
symbol	symbols or numbers	pattern	?, 1, ..., ??!?, /	40	3181
slang	DotA-specific game-technical terms, used to coordinate with team	list	ursa, mid, back, farm, bkb	30	10046
command	in-game commands, control words to trigger certain effects	pattern	!ff, !pause, -swap	20	2513
stop	English stop words	list	was, i, it, can, you	10	1322
timemark	automatically generated time-stamps, prepended in pause-mode	pattern	[00:05], [01:23]	5	223

Fig. 2. Average use of annotated words per chat-mode. Chat-mode depicted as solid bars (all-chat) and as transparent bars (ally-chat). Error-markers show one standard deviation. Category "unannotated" was omitted.

players in the same team as the sender. We observe that on average 90% of all messages are exchanged in the ally-chat and only 10% are broadcasted to all players. Private player-to-player communication is also possible, but not saved within our data. Figure 2 shows the relative amount of annotated words averaged over all matches for both chat modes. It is interesting to see that words from the "stop"-category are used almost equally in both chat modes, meaning that our selection of stop-words is context-independent. The usage of words from the "slang"-category is twice as high in the ally-chat, since slang is mainly used to transfer sensitive information to coordinate the team in its battle. The heavy relative use of slang in the ally-chat creates a bias in almost all other annotation categories towards the all-chat.

C. Toxicity detection

For the purpose of our investigation, we define toxicity as the use of profane language by one player to insult or humiliate a different player *in his own team*. As such, the use of "bad" words is a necessary, but not sufficient condition for toxicity. For example, bad words can also be used just to curse without the intent to actually insult someone else. Profanity is also used in ironic or humoristic ways. For example, some players use self-deprecating remarks to admit in-game mistakes: "sry, I am such a noob - lol". Thus, detection of toxicity can not be based on words alone but needs to take the current context into account.

We are using n-grams to distinguish toxicity from ordinary profane language. An n-gram is a contiguous sequence of n words that appears in a context. The context in our case consists of all words in the chat-line that contained the "bad" word plus all words from all chat-lines that were sent by the same player to the ally-chat not more than 1 second before or after.

For all players who participated in at least 10 matches, we search for all "bad" words they use, construct their contexts and count each n-gram that contains at least one "bad" word for $n = 1, 2, 3, 4$. Afterwards, we look at the 100 most frequently used n-grams for $n = 1, 2, 3, 4$ and manually determine which of them are toxic and which are not. Our criterion for toxicity is the following: for unigrams ($n = 1$) we consider them toxic if they could be understood as an insult. For example "crap" is no insult, but "moron" is. For n-grams with $n = 2, 3, 4$, we consider every context toxic that includes an insult directed towards a person. Examples include "f*ing idiot", "shut the f*" and "i hope u die". On the contrary, profane language that we do not classify as toxic includes n-grams like "f* this","cant do s*" and "dont give a f*".[7] In total, we deem 45 unigrams, 21 bigrams, 32 trigams and 36 quadgrams as toxic. The list of these n-grams is provided as supplement to this work online (see Footnote 6).

IV. ANALYSIS OF GAME TOXICITY AND SUCCESS

A. Triggers of toxicity

Our method detects at least one toxic remark in 6528 out of the 10305 matches. In 90% of all toxic matches, there are at most 5 toxic remarks detected. Several outliers exist in the data, the strongest contains 22 toxic remarks in a single match. The total number of toxic remarks was 16950. We expect that certain game events trigger players to act toxic.

[7]Some authors did not want to include explicit quotes of profanity, which is why we decided to apply self-censorship in this work.

Fig. 3. Correlogram between annotation categories and winrate.

	kill-events	from killer	from victim
toxicity ($\Delta = 5s$)	2219	23	849
random ($\Delta = 5s$)	1488	74	478
toxicity ($\Delta = 10s$)	5285	124	2559
random ($\Delta = 10s$)	3176	200	1042

One possible game event is a *kill* where one player (killer) temporarily eliminates the character of another player in the opposing team (victim). There is a reaction time Δ involved between the actual kill-event and the time a player needs to submit a response to the chat. We look for each toxic remark if there was a kill-event taking place not earlier than Δ seconds before. For comparison, we also choose 16950 random chat-lines (distributed over all matches) and look for a kill-event in their recent past as well. It turns out, that **toxic remarks are more frequently preceded by kill-events than random remarks**. Table II reports the absolute number of kill-events and how many of them were submitted by the killer or the victim. Especially victims of kill-events tend to become toxic, potentially blaming their teammates for their own fate.

B. Game success and profanity

We have the hypothesis that with diminishing chances to succeed in the game, the level of profanity raises. To test our hypothesis, we compute the winrate for each player as the amount of matches won divided by the amount of matches played in total. We restrict the analysis to players who participated in at least 10 matches, which leaves 4009 distinct players in our dataset. Next, we count how many words the players used for our annotation categories "bad", "praise" and "slang". Normalized by the total number of words, we correlate this number with the winrate, and plot the results in Figure 3. Surprisingly, there seems to be no strong linear correlation in either case, which is confirmed by the correlation-matrix given by Table III.

An analysis based on absolute word-counts with focus on whole teams (rather than single players) reveals a different picture: for each "bad" word used by a winning team, we determine the point of time in the match when it was submitted to the chat. As different matches vary in duration (recall Figure 1) we normalize time to the interval $[0, 1]$ on the
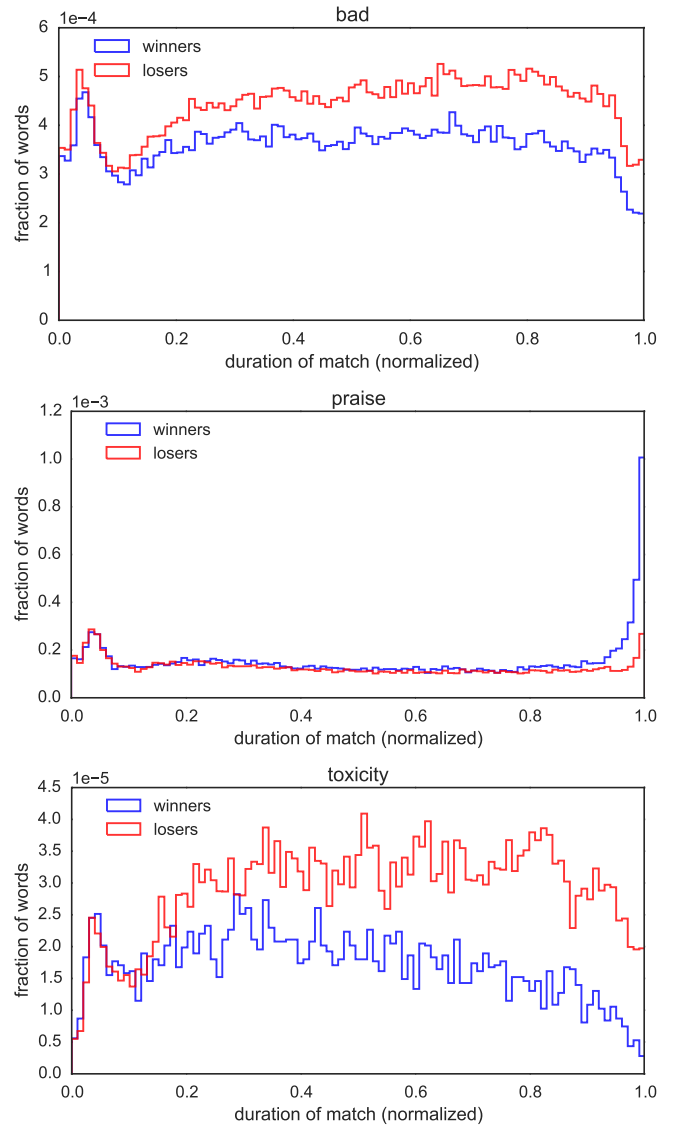

Fig. 4. Overlapping histograms, comparing winning and losing teams in their usage of words from categories "bad", "praise", and toxic n-grams.

horizontal axis, with 1 indicating the end of a match. Out of this data we construct a histogram using 100 equally distributed bins. We overlap this histogram with a second histogram, constructed the same way but for words used by the losing teams. As winning and losing teams use a slightly different absolute number of words per bin, we normalize each bin accordingly to eliminate bias. Figure 4 reports on the vertical axis the fraction of words used in each bin over all words used by the respective winning or losing team. It also shows a histogram based on toxicity detected by our toxic n-grams.

As we can see from the top part of Figure 4, after a short initial period, in which it is uncertain to the players whether they might be losing or winning, we observe that teams that will lose the match in the end tend to use relatively more bad words than teams that will win the match. This difference is

	winrate	bad	praise	slang
winrate	1.0	0.0739	-0.0161	0.0059
bad	-0.0739	1.0	0.0454	-0.1540
praise	-0.0161	0.0454	1.0	0.1152
slang	0.0059	-0.1540	0.1152	1.0

Fig. 5. Performance of the linear SVM on ally-chats.

even bigger if toxicity is considered. More interestingly: while the usage of bad words is somewhat consistent throughout the match, the usage of toxicity varies more. It seems that **the winning teams use less toxicity at the late stages of the match**, as it becomes apparent that they will be victorious. **The need to shame and blame teammates seems to be significantly higher for the losing team than the winning team at this point in time**. Another interesting aspect is the usage of the category "praise" which seems consistent for most of the matches but peaks clearly for the winning team by the very end. This effect is due to the traditional phrase "gg" (good game) which is a word from the "praise" category and often used just before the match finishes. Winning teams use this phrase significantly more, probably as they might perceive the match as more enjoyable.

C. Predicting match outcome

As we have shown, toxicity appears only in 60% of all matches and is thus too infrequent to be used for predicting match outcome in general. Therefore, we will analyse the predictive power of all words with respect to their annotations, including the category of "bad" words. We train a linear support vector machine (SVM) to predict the winning team on a feature-set based on TF-IDF (term frequency inverse document frequency) of each word, a standard weighting technique frequently used in information retrieval [2]. For all computations, we use Scikit-learn [3] with its default parameters for all algorithms and do not undertake any effort to optimize them. The idea is *not* to create the most accurate classifier possible but rather to use the accuracy of the classifier to measure the importance of words with respect to match outcome.

The outcome should become more certain with the progression of the match, which should be reflected in the words used by the players. We introduce the parameter t to control the amount of chat history that is given to the classifier. For example, for $t = 1.0$ the classifier is trained (and evaluated) on the complete ally-chats of each match, whereas for $t = 0.5$ it only knows what was written until the middle of the matches. The classifier itself has no notion of time: the TF-IDF features are purely based on frequencies and reflect neither order of words nor the specific time they were submitted to the chat.

As each word corresponds to one feature, we can partition all features by using our annotation system. We use the classifier 1) for all words regardless of their annotations, 2) for all words but words from the "command" category, 3) for

no words except from the "bad" category and 4) for no words except from the "slang" category. The reason for excluding words from "command" is to avoid to provide the classifier with information if a player forfeited, which is announced by typing the command word "!ff" in the chat.

Figure 5 shows the average accuracy and the 95% confidence interval of the classifier for these scenarios under a 10-fold cross-validation. The number of used features and the accuracy scores for $t = 0.5$, $t = 0.75$ and $t = 1.0$ are presented in Table IV.

While words from the "bad" category (which constitute a precondition for toxicity) have some predictive power, it is significantly lower than using just *all words* or words from "slang" alone. We find it also remarkable that "slang" uses the least amount of features but gives still fairly good predictions. This might be due to the importance of team coordination which is covered mostly by key words from this category. It seems reasonable that their usage shows not only the game expertise of players, but also engagement and an increased interest to improve the team-play, which could result in a better chance to win the match. The occurrence of "bad" words however seems to be much less indicative for either winning or losing, suggesting only a weak link to game success. Consequently, also toxicity might not be the best indicator to determine if a game is going well for a team or not. Profanity will appear either way.

V. RELATED WORK

Antisocial behavior in virtual environments has been investigated in the field of computer sociology, most commonly under the term cyberbullying [4]. The impact of profane language on video games [5] and in a wider sense also on social media [6] is a vital area of research. Suler [7] shows psychological factors explaining the *online disinhibition effect*, giving *toxic disinhibition* as a negative example. This effect is a possible explanation why we observe such high levels of bad behavior online in general.

Similar to toxicity is the concept of *griefing*, the act of disrupting the game experience of other players by performing unacceptable actions. This has been investigated for virtual

TABLE IV
PERFORMANCE OF CLASSIFIER

	t = 0.5			t = 0.75			t = 1.0		
	#features	avg accuracy	std accuracy	#features	avg accuracy	std accuracy	#features	avg accuracy	std accuracy
all words	127612	0.6399	0.0140	170063	0.7689	0.0092	208598	0.9407	0.0048
all but "command"	126900	0.6346	0.0103	169298	0.7421	0.0099	207758	0.8708	0.0070
only "bad"	1442	0.5720	0.0137	1767	0.6077	0.0096	2020	0.6538	0.0108
only "slang"	880	0.5877	0.0189	908	0.6875	0.0114	921	0.8295	0.0093

worlds like Second Life [8] and MMORPGs like World of Warcraft [9].

An excellent case study for toxicity in MOBAs is given by the works of Blackburn and Kwak [10], [11]. The authors use crowd-sourced decisions from the *tribunal*, a player-based court that passes judgement on reported incidents in matches from League of Legends.[8] While our definition of toxicity is tied to profane language only, the authors additionally consider certain in-game actions (i.e. "intentional feeding") as toxic. They develop a classifier to assists or even substitute the crowd-sourced decisions of the tribunal, which are whether an accused player is guilty of toxic behavior or not. As only cases submitted to the tribunal are considered, the authors have access to a ground truth for toxicity which is not present for our data. However, this might also create a selection bias, as typical matches will not end up on the tribunal.

Shim et al. [12] describe a different system based on the Pagerank to filter out "bad players" in MOBAs. Our approach is orthogonal, as it uses natural language processing on the player chats instead of relying on player's complaints submitted via a report function. Institutions like the tribunal would not work without players reporting others, while our approach does not need any explicit player feedback to detect and monitor toxicity.

VI. DISCUSSION AND CONCLUSION

We developed a methodology to annotate frequently used expressions in written chat communication of Multiplayer Online Games. While our method is tested in this work only with data from DotA, we believe that it can be adapted to other MOBAs and possibly even games of different genres. To use the full system, one would need to update the pre-defined lists and patterns of game-specific terms to match their equivalents from the new game. Although this requires some degree of game-knowledge, the detection of profanity itself is largely independent from any game-specifics, as it is based on profanity used in English language, enriched by a few terms commonly used in computer games.

The developed toxicity detection is based on contextual information to distinguish simple swearing from deliberate insults. It can be a building block for a monitoring system that can be used together with player reports to identify toxic players. Moreover, a well-trained classifier could be used to

design a live-system that displays the odds of winning for each team to observers based on their communication, while the match is still ongoing.

Our analysis shows that toxicity is fueled by the inherent competitiveness (i.e., killing each other) of MOBA games but is only weakly linked to success. If players can be successful despite being toxic, they need a different incentive to cease insulting and behave more pleasant. On the other hand, the matchmaking systems that ensemble the teams could be altered to take toxicity into account to avoid creating a social powder keg. Although we might not be able to prevent toxicity entirely, controlling it would ensure a much more positive game experience for newcomers and experienced players alike.

REFERENCES

[1] F. Kuipers, R. Kooij, D. De Vleeschauwer, and K. Brunnström, "Techniques for measuring quality of experience," in *Wired/wireless internet communications*. Springer, 2010, pp. 216–227.

[2] C. D. Manning, P. Raghavan, H. Schütze *et al.*, *Introduction to information retrieval*. Cambridge university press Cambridge, 2008, vol. 1.

[3] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay, "Scikit-learn: Machine learning in Python," *Journal of Machine Learning Research*, vol. 12, pp. 2825–2830, 2011.

[4] P. K. Smith, J. Mahdavi, M. Carvalho, S. Fisher, S. Russell, and N. Tippett, "Cyberbullying: Its nature and impact in secondary school pupils," *Journal of child psychology and psychiatry*, vol. 49, no. 4, pp. 376–385, 2008.

[5] A. H. Ivory and C. E. Kaestle, "The effects of profanity in violent video games on players' hostile expectations, aggressive thoughts and feelings, and other responses," *Journal of Broadcasting & Electronic Media*, vol. 57, no. 2, pp. 224–241, 2013.

[6] S. Sood, J. Antin, and E. Churchill, "Profanity use in online communities," in *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*. ACM, 2012, pp. 1481–1490.

[7] J. Suler, "The online disinhibition effect," *Cyberpsychology & behavior*, vol. 7, no. 3, pp. 321–326, 2004.

[8] T. Chesney, I. Coyne, B. Logan, and N. Madden, "Griefing in virtual worlds: causes, casualties and coping strategies," *Information Systems Journal*, vol. 19, no. 6, pp. 525–548, 2009.

[9] C. Y. Foo and E. M. Koivisto, "Defining grief play in mmorpgs: player and developer perceptions," in *Proceedings of the 2004 ACM SIGCHI International Conference on Advances in computer entertainment technology*. ACM, 2004, pp. 245–250.

[10] J. Blackburn and H. Kwak, "STFU NOOB!: predicting crowdsourced decisions on toxic behavior in online games," in *Proceedings of the 23rd international conference on World wide web*. ACM, 2014, pp. 877–888.

[11] H. Kwak and J. Blackburn, "Linguistic analysis of toxic behavior in an online video game," in *Social Informatics*. Springer, 2014, pp. 209–217.

[12] J. Y. Shim, T. H. Kim, and S. W. Kim, "Decision support of bad player identification in moba games using pagerank based evidence accumulation and normal distribution based confidence interval," *International Journal of Multimedia & Ubiquitous Engineering*, vol. 9, no. 8, 2014.

[8]By the time of this work, the tribunal together with its data has not been accessible for over a year as it undergoes maintenance by Riot Games.